

sign_language_cnn_100_percent

May 5, 2020

```
In [2]: import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import pandas as pd
from tensorflow import keras
from tensorflow.keras.layers import Conv2D, Flatten, MaxPooling2D, AveragePooling2D, Dense
from keras.preprocessing.image import ImageDataGenerator
import string

from libitmal import kernelfuns as itmalkernelfuns
itmalkernelfuns.EnableGPU()

%matplotlib inline

def get_mnist_dataset():

    test_pd = pd.read_csv("./SignLanguageData/sign_mnist_test.csv",
                           skiprows=1)
    train_pd = pd.read_csv("./SignLanguageData/sign_mnist_train.csv",
                            skiprows=1)

    return train_pd, test_pd

train_pd, test_pd = get_mnist_dataset()
X_train, X_test = train_pd.values[:,1:], test_pd.values[:,1:]
y_train, y_test = train_pd.values[:,0], test_pd.values[:,0]

class_names = list(string.ascii_lowercase)

train_pd.head()
```

Out [2]:

3	107	118	127	134	139	143	146	150	153	...	207.4	207.5	207.6	\
0	6	155	157	156	156	157	156	158	158	...	69	149	128	
1	2	187	188	188	187	187	186	187	188	...	202	201	200	
2	2	211	211	212	212	211	210	211	210	...	235	234	233	
3	13	164	167	170	172	176	179	180	184	...	92	105	105	
4	16	161	168	172	173	178	184	189	193	...	76	74	68	

	207.7	206.4	206.5	206.6	204.6	203.8	202.13
0	87	94	163	175	103	135	149
1	199	198	199	198	195	194	195
2	231	230	226	225	222	229	163
3	108	133	163	157	163	164	179
4	62	53	55	48	238	255	255

[5 rows x 785 columns]

```
In [3]: X_train = X_train / 255
X_test = X_test / 255
```

```
X_train = X_train.reshape(*X_train.shape[:1], 28, 28)
X_test = X_test.reshape(*X_test.shape[:1], 28, 28)
```

```
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
```

```
batch_size, height, width, channel = X_train.shape
```

```
print(X_train.shape)
```

```
(27454, 28, 28, 1)
```

```
In [185]: # Datageneration form https://www.kaggle.com/madz2000/cnn-using-keras-99-7-accuracy
datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range=10, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.1, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total)
    horizontal_flip=False, # randomly flip images
    vertical_flip=False # randomly flip images
)
```

```
datagen.fit(X_train)
```

```
In [6]: def create_le_net():
    model = keras.models.Sequential([
        ZeroPadding2D(input_shape=X_train.shape[1:], padding=(3, 3)),
        Conv2D(filters=6, kernel_size=(5, 5), strides=1, activation="tanh"),
        AveragePooling2D(pool_size=6, strides=2, padding="same"),
```

```

        Conv2D(filters=16, kernel_size=(5, 5), strides=1, activation="tanh"),
        AveragePooling2D(pool_size=6, strides=2, padding="same"),
        Conv2D(filters=120, kernel_size=(5, 5), strides=1, activation="tanh"),
        Flatten(),
        Dense(84, activation="tanh"),
        Dense(len(class_names), activation="softmax")
    ])

    return model

def create_model():
    model = keras.models.Sequential([
        ZeroPadding2D(input_shape=X_train.shape[1:], padding=(3, 3)),
        Conv2D(filters=16, kernel_size=(3, 3), activation="relu"),
        AveragePooling2D(),
        Conv2D(filters=32, kernel_size=(3, 3), activation="relu"),
        AveragePooling2D(),
        Flatten(),
        Dense(256, activation="relu"),
        Dense(512, activation="relu"),
        Dense(128, activation="relu"),
        Dense(64, activation="relu"),
        Dense(len(class_names), activation="softmax")
    ])

    return model

```

```
In [201]: X_test.shape
```

```
Out[201]: (7171, 28, 28, 1)
```

```
In [187]: # LeNet
```

```
model = create_model()
```

```
In [188]: #keras.utils.plot_model(model, "my_mnist_model.png", show_shapes=True)
```

```
In [189]: model.compile(loss="sparse_categorical_crossentropy",
                        optimizer="adam",
                        metrics=["accuracy"])
```

```
In [190]: from keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
early_stopping = EarlyStopping(monitor='loss',
                               patience=30,
                               verbose=0,
                               mode='min')
```

```
mcp_save = ModelCheckpoint('cnn_model_checkpoint.h5',
                           save_best_only=True,
```

```

        monitor='val_loss',
        mode='min')

    history = model.fit(datagen.flow(X_train,
                                    y_train,
                                    batch_size = 128,
                                    validation_split=0.2),
                        epochs=500,
                        validation_split=0.2,
                        callbacks=[early_stopping, mcp_save])

Epoch 1/500
215/215 [=====] - 5s 22ms/step - loss: 2.6089 - acc: 0.1943
Epoch 2/500
215/215 [=====] - 4s 19ms/step - loss: 1.4556 - acc: 0.5055
Epoch 3/500
215/215 [=====] - 4s 20ms/step - loss: 0.8565 - acc: 0.7064
Epoch 4/500
215/215 [=====] - 4s 19ms/step - loss: 0.5162 - acc: 0.8204
Epoch 5/500
215/215 [=====] - 4s 20ms/step - loss: 0.3500 - acc: 0.8827
Epoch 6/500
215/215 [=====] - 4s 20ms/step - loss: 0.2634 - acc: 0.9108
Epoch 7/500
215/215 [=====] - 4s 20ms/step - loss: 0.1906 - acc: 0.9372
Epoch 8/500
215/215 [=====] - 4s 20ms/step - loss: 0.1396 - acc: 0.9549
Epoch 9/500
215/215 [=====] - 4s 20ms/step - loss: 0.1299 - acc: 0.9565
Epoch 10/500
215/215 [=====] - 4s 20ms/step - loss: 0.1100 - acc: 0.9629
Epoch 11/500
215/215 [=====] - 4s 20ms/step - loss: 0.0986 - acc: 0.9682
Epoch 12/500
215/215 [=====] - 4s 19ms/step - loss: 0.0891 - acc: 0.9703
Epoch 13/500
215/215 [=====] - 4s 19ms/step - loss: 0.0690 - acc: 0.9774
Epoch 14/500
215/215 [=====] - 4s 19ms/step - loss: 0.0637 - acc: 0.9788
Epoch 15/500
215/215 [=====] - 4s 20ms/step - loss: 0.0568 - acc: 0.9819
Epoch 16/500
215/215 [=====] - 4s 19ms/step - loss: 0.0543 - acc: 0.9820
Epoch 17/500
215/215 [=====] - 4s 20ms/step - loss: 0.0496 - acc: 0.9832
Epoch 18/500
215/215 [=====] - 4s 19ms/step - loss: 0.0553 - acc: 0.9820
Epoch 19/500

```

215/215 [=====] - 4s 20ms/step - loss: 0.0540 - acc: 0.9825
 Epoch 20/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0378 - acc: 0.9880
 Epoch 21/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0424 - acc: 0.9862
 Epoch 22/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0346 - acc: 0.9890
 Epoch 23/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0317 - acc: 0.9896
 Epoch 24/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0388 - acc: 0.9872
 Epoch 25/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0342 - acc: 0.9889
 Epoch 26/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0291 - acc: 0.9902
 Epoch 27/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0288 - acc: 0.9909
 Epoch 28/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0269 - acc: 0.9907
 Epoch 29/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0291 - acc: 0.9902
 Epoch 30/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0250 - acc: 0.9924
 Epoch 31/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0297 - acc: 0.9907
 Epoch 32/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0252 - acc: 0.9920
 Epoch 33/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0292 - acc: 0.9910
 Epoch 34/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0274 - acc: 0.9913
 Epoch 35/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0257 - acc: 0.9916
 Epoch 36/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0188 - acc: 0.9941
 Epoch 37/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0178 - acc: 0.9943
 Epoch 38/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0195 - acc: 0.9943
 Epoch 39/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0242 - acc: 0.9924
 Epoch 40/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0182 - acc: 0.9942
 Epoch 41/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0204 - acc: 0.9931
 Epoch 42/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0194 - acc: 0.9939
 Epoch 43/500

215/215 [=====] - 4s 19ms/step - loss: 0.0181 - acc: 0.9940
 Epoch 44/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0215 - acc: 0.9927
 Epoch 45/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0146 - acc: 0.9956
 Epoch 46/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0201 - acc: 0.9932
 Epoch 47/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0205 - acc: 0.9939
 Epoch 48/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0153 - acc: 0.9947
 Epoch 49/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0164 - acc: 0.9948
 Epoch 50/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0140 - acc: 0.9962
 Epoch 51/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0134 - acc: 0.9957
 Epoch 52/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0177 - acc: 0.9944
 Epoch 53/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0165 - acc: 0.9946
 Epoch 54/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0212 - acc: 0.9937
 Epoch 55/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0111 - acc: 0.9964
 Epoch 56/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0148 - acc: 0.9956
 Epoch 57/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0177 - acc: 0.9944
 Epoch 58/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0153 - acc: 0.9956
 Epoch 59/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0062 - acc: 0.9982
 Epoch 60/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0112 - acc: 0.9964
 Epoch 61/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0269 - acc: 0.9916
 Epoch 62/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0085 - acc: 0.9972
 Epoch 63/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0167 - acc: 0.9948
 Epoch 64/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0083 - acc: 0.9972
 Epoch 65/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0125 - acc: 0.9962
 Epoch 66/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0100 - acc: 0.9969
 Epoch 67/500

215/215 [=====] - 4s 20ms/step - loss: 0.0125 - acc: 0.9958
 Epoch 68/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0127 - acc: 0.9960
 Epoch 69/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0160 - acc: 0.9949
 Epoch 70/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0116 - acc: 0.9964
 Epoch 71/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0149 - acc: 0.9956
 Epoch 72/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0094 - acc: 0.9970
 Epoch 73/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0098 - acc: 0.9970
 Epoch 74/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0062 - acc: 0.9983
 Epoch 75/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0084 - acc: 0.9974
 Epoch 76/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0077 - acc: 0.9974
 Epoch 77/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0165 - acc: 0.9954
 Epoch 78/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0093 - acc: 0.9971
 Epoch 79/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0156 - acc: 0.9955
 Epoch 80/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0058 - acc: 0.9984
 Epoch 81/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0080 - acc: 0.9980
 Epoch 82/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0134 - acc: 0.9955
 Epoch 83/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0119 - acc: 0.9967
 Epoch 84/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0130 - acc: 0.9960
 Epoch 85/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0047 - acc: 0.9985
 Epoch 86/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0137 - acc: 0.9957
 Epoch 87/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0104 - acc: 0.9966
 Epoch 88/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0125 - acc: 0.9957
 Epoch 89/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0045 - acc: 0.9984
 Epoch 90/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0035 - acc: 0.9988
 Epoch 91/500

215/215 [=====] - 4s 19ms/step - loss: 0.0152 - acc: 0.9951
 Epoch 92/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0080 - acc: 0.9977
 Epoch 93/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0105 - acc: 0.9969
 Epoch 94/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0134 - acc: 0.9966
 Epoch 95/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0112 - acc: 0.9965
 Epoch 96/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0050 - acc: 0.9984
 Epoch 97/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0084 - acc: 0.9969
 Epoch 98/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0129 - acc: 0.9965
 Epoch 99/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0088 - acc: 0.9977
 Epoch 100/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0074 - acc: 0.9977
 Epoch 101/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0070 - acc: 0.9979
 Epoch 102/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0052 - acc: 0.9985
 Epoch 103/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0059 - acc: 0.9981
 Epoch 104/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0144 - acc: 0.9957
 Epoch 105/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0058 - acc: 0.9985
 Epoch 106/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0078 - acc: 0.9977
 Epoch 107/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0085 - acc: 0.9971
 Epoch 108/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0105 - acc: 0.9968
 Epoch 109/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0058 - acc: 0.9983
 Epoch 110/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0053 - acc: 0.9985
 Epoch 111/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0072 - acc: 0.9979
 Epoch 112/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0085 - acc: 0.9973
 Epoch 113/500
 215/215 [=====] - 4s 19ms/step - loss: 0.0094 - acc: 0.9974
 Epoch 114/500
 215/215 [=====] - 4s 20ms/step - loss: 0.0091 - acc: 0.9973
 Epoch 115/500


```

215/215 [=====] - 4s 20ms/step - loss: 0.0044 - acc: 0.9988
Epoch 116/500
215/215 [=====] - 4s 19ms/step - loss: 0.0095 - acc: 0.9972
Epoch 117/500
215/215 [=====] - 4s 20ms/step - loss: 0.0049 - acc: 0.9983
Epoch 118/500
215/215 [=====] - 4s 20ms/step - loss: 0.0106 - acc: 0.9967
Epoch 119/500
215/215 [=====] - 4s 20ms/step - loss: 0.0080 - acc: 0.9977
Epoch 120/500
215/215 [=====] - 4s 19ms/step - loss: 0.0036 - acc: 0.9988

```

```
In [191]: import pandas as pd
```

```

evaluation = model.evaluate(X_test, y_test)
print(f'Model evaluation: {evaluation}')


pd.DataFrame(history.history).plot(figsize=(14, 10))
plt.grid(True)
plt.gca().set_ylim(-0.1, 1.1)
plt.xlabel("Epoch [N]")
plt.ylabel("Percentage [%]")
plt.title("Loss and Accuracy as a function of epochs")
plt.show()

```

```

7171/7171 [=====] - 1s 71us/step
Model evaluation: [0.0008896882954271737, 1.0]

```



```
In [192]: import math
```

```
num_rows = 3
num_cols = 3
```

```
X_new = X_test[:num_rows*num_cols]
```

```
y_pred = model.predict_classes(X_new)
```

```
fig, ax = plt.subplots(num_rows, num_cols, figsize=(18, 16))
```

```
for index, image in enumerate(X_new):
```

```
    ax[math.floor(index/num_rows), index%num_rows].imshow(image.reshape((28,28)))
```

```
    ax[math.floor(index/num_rows), index%num_rows].set_title(
```

```
        f"Actual: {class_names[y_test[index]]}\nPredicted: {class_names[y_pred[index]]"
```

```
        fontsize=16)
```

```
    ax[int(index/num_rows), index%num_rows].axis('off')
```

```
fig.tight_layout()
```

```
fig.suptitle(f'First Predictions', fontsize=20)
```

```
fig.subplots_adjust(top=0.88)
```

```
fig.show()
```

First Predictions



```
In [193]: y_pred = model.predict_classes(X_test)

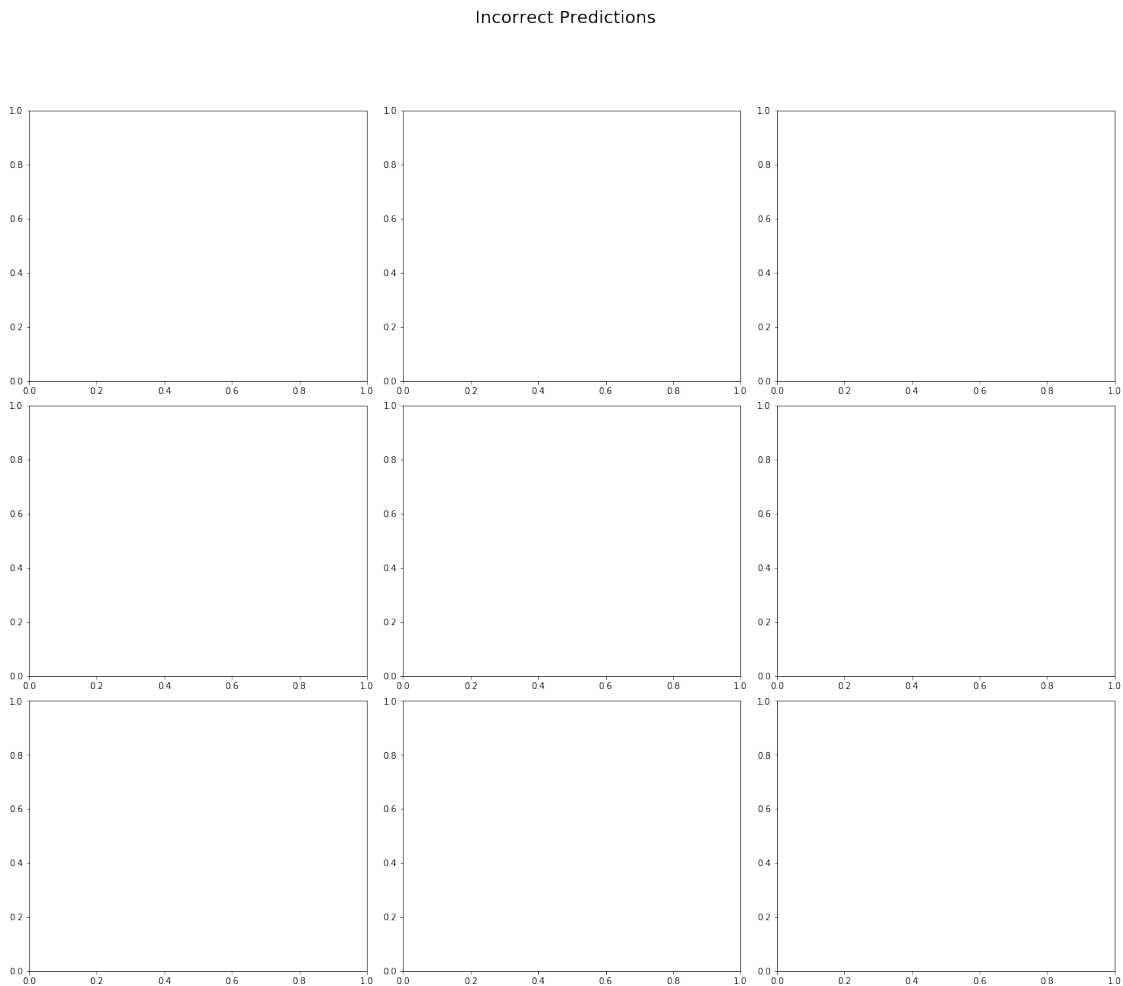
confusion_indices = np.where(y_pred != y_test)
X_confusion = X_test[confusion_indices]
y_pred_confusion = y_pred[confusion_indices]
y_test_confusion = y_test[confusion_indices]

fig, ax = plt.subplots(num_rows, num_cols, figsize=(18, 16))
for index, image in enumerate(X_confusion[:num_rows*num_cols]):
    ax[math.floor(index/num_rows), index%num_rows].imshow(image.reshape((28,28)), inter=
    ax[math.floor(index/num_rows), index%num_rows].set_title(
        f"Actual: {class_names[y_test_confusion[index]]}\nPredicted: {class_names[y_pr
        fontsize=16)
    ax[int(index/num_rows), index%num_rows].axis('off')
```

```

fig.tight_layout()
fig.suptitle('Incorrect Predictions', fontsize=20)
fig.subplots_adjust(top=0.88)
fig.show()

```



```
In [194]: import sklearn.metrics as metrics
```

```

confusion_matrix = metrics.confusion_matrix(y_test, y_pred)
row_sum = confusion_matrix.sum(axis=1, keepdims=True)
norm_confusion_matrix = confusion_matrix / row_sum

```

```

# Because j and z aren't possible we cant include them in confusion matrix
class_names_clean = class_names.copy()
class_names_clean.remove('j')
class_names_clean.remove('z')

```

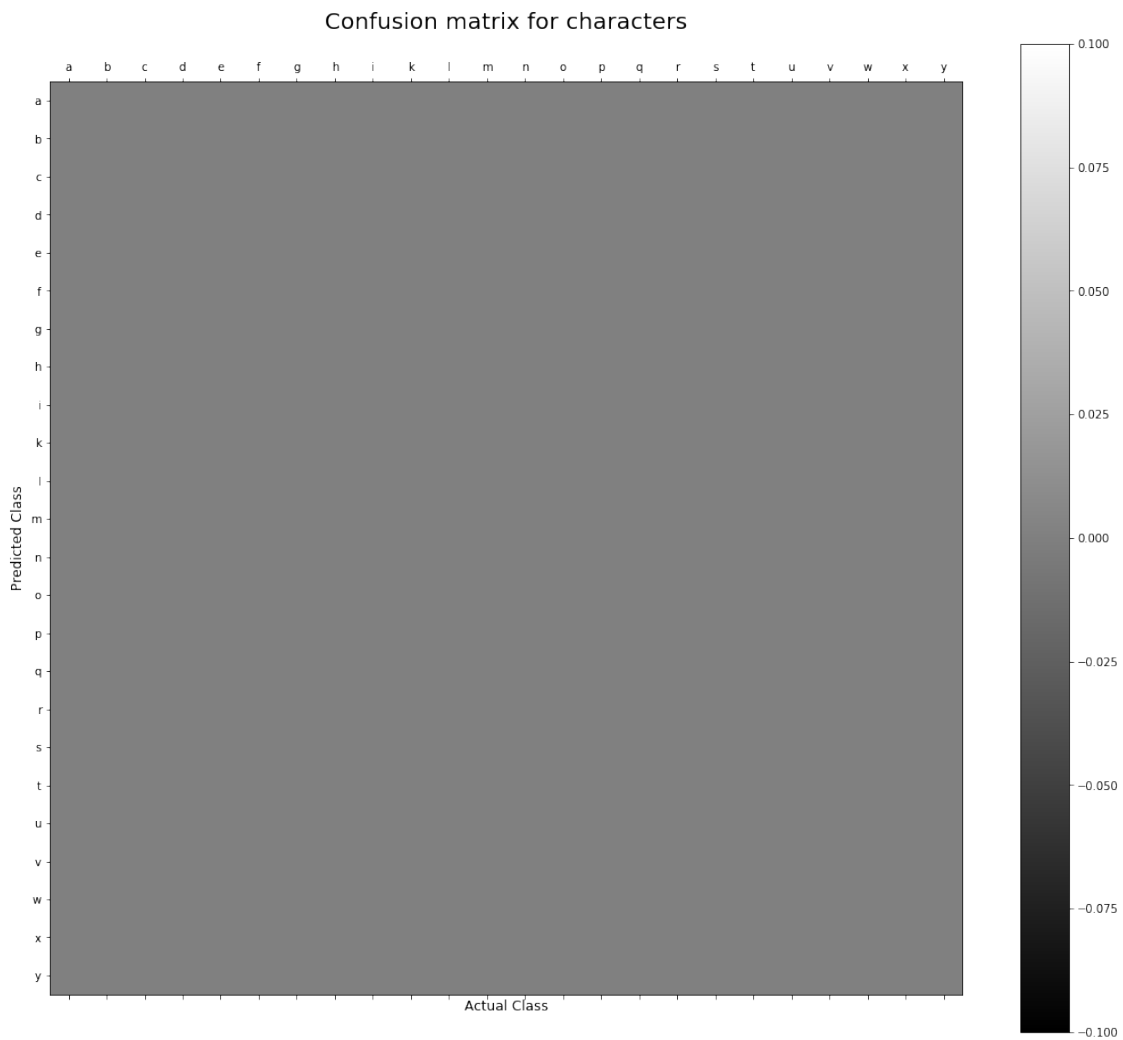
```

np.fill_diagonal(norm_confusion_matrix, 0)

fig, ax = plt.subplots(figsize=(18, 16))

mat_ax = ax.matshow(norm_confusion_matrix, interpolation='nearest', cmap=plt.cm.gray)
fig.colorbar(mat_ax)
ax.set_title('Confusion matrix for characters', fontsize=20)
ax.set_xlabel('Actual Class', fontsize=12)
ax.set_ylabel('Predicted Class', fontsize=12)
ax.set_xticks(ticks=np.arange(0, len(class_names_clean)))
ax.set_xticklabels(class_names_clean)
ax.set_yticks(ticks=np.arange(0, len(class_names_clean)))
ax.set_yticklabels(class_names_clean)
fig.show()

```



```
In [195]: y_pred = model.predict_classes(X_test)
          print(metrics.f1_score(y_test, y_pred, average="micro"))
```

1.0

```
In [196]: model.save('cnn_model_100.h5')
```