

LOG_SignLanguage

May 2, 2020

```
In [7]: from sklearn.pipeline import Pipeline
        from sklearn.preprocessing import StandardScaler
        from sklearn.svm import SVC
        import pandas as pd
        import numpy as np
        from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import LabelEncoder
        from sklearn.decomposition import PCA
        import time
        import matplotlib.pyplot as plt

train = pd.read_csv('./SignLanguageData/sign_mnist_train.csv')
test = pd.read_csv('./SignLanguageData/sign_mnist_test.csv')

train = train.astype(float)
test = test.astype(float)

Y_train = train['label'].values
train.drop('label', axis = 1, inplace = True)

Y_test = test['label'].values
test.drop('label', axis = 1, inplace = True)

labelizer = LabelEncoder()
Y_train = labelizer.fit_transform(Y_train)
Y_test = labelizer.transform(Y_test)

X_train = train.values
X_test = test.values

pipeline = Pipeline([
    ("std_scale", StandardScaler()),
    ("pca", PCA(.95)),
])

X_train_transform = pipeline.fit_transform(X_train, Y_train)
```

```


X_test_transform = pipeline.transform(X_test)
plt.imshow(X_train[0].reshape(28,28))
print(X_test.shape)
print(X_train.shape)
print(Y_train[0])
print(Y_test[0])

```

```

(7172, 784)
(27455, 784)
3
6

```



```

In [2]: # This class is taken from lesson 08, Machine Learning.
from sklearn.metrics import classification_report, f1_score
from time import time
from sklearn.model_selection import GridSearchCV

def SearchReport(model):

    def GetBestModelCTOR(model, best_params):
        def GetParams(best_params):
            ret_str=""
            for key in sorted(best_params):
                value = best_params[key]
                temp_str = "" if str(type(value))=="<class 'str'>" else ""

```

```

        if len(ret_str)>0:
            ret_str += ','
            ret_str += f'{key}={temp_str}{value}{temp_str}'
        return ret_str
    try:
        param_str = GetParams(best_params)
        return type(model).__name__ + '(' + param_str + ')'
    except:
        return "N/A(1)"

print("\nBest model set found on train set:")
print()
print(f"\tbest parameters={model.best_params}")
print(f"\tbest '{model.scoring}' score={model.best_score}")
print(f"\tbest index={model.best_index}")
print()
print(f"Best estimator C TOR:")
print(f"\t{model.best_estimator}")
print()
try:
    print(f"Grid scores ('{model.scoring}') on development set:")
    means = model.cv_results_['mean_test_score']
    stds = model.cv_results_['std_test_score']
    i=0
    for mean, std, params in zip(means, stds, model.cv_results_['params']):
        print("\t[%2d]: %0.3f (+/-%0.03f) for %r" % (i, mean, std * 2, params))
        i += 1
except:
    print("WARNING: the random search do not provide means/stds")

assert "f1_micro"==str(model.scoring), f"come on, we need to fix the scoring to be a
return f"best: data=SignLanguageMnist, score={model.best_score:.05f}, model={GetBes

def ClassificationReport(model, X_test, y_test, target_names=None):
    assert X_test.shape[0]==y_test.shape[0]
    print("\nDetailed classification report:")
    print("\tThe model is trained on the full development set.")
    print("\tThe scores are computed on the full evaluation set.")
    print()
    y_true, y_pred = y_test, model.predict(X_test)
    print(classification_report(y_true, y_pred, target_names))
    print()

def FullReport(model, X_test, y_test, t):
    print(f"SEARCH TIME: {t:0.2f} sec")
    beststr, bestmodel = SearchReport(model)
    ClassificationReport(model, X_test, y_test)
    print(f"CTOR for best model: {bestmodel}\n")

```

```

        print(f"{beststr}\n")
        return beststr, bestmodel

In [ ]: from sklearn.linear_model import LogisticRegression
        from sklearn.preprocessing import StandardScaler
        from sklearn.pipeline import Pipeline
        from sklearn.decomposition import PCA

        import sys
        old_stdout = sys.stdout
        sys.stdout = open('output_log.txt', 'w')

        model_pipe = Pipeline([
            ("std_scaler", StandardScaler())
        ])

        tuning_parameters_1 = {
            'solver':('newton-cg', 'lbfgs', 'sag'),
            'max_iter':[100, 1000, 10000],
            'C': [1, 10, 100]
        }

        tuning_parameters_2 = {
            'solver':('liblinear', 'saga'),
            'penalty':('l1', 'l2'),
            'max_iter':[100, 1000, 10000],
            'C': [1, 10, 100]
        }

        tuning_array = [tuning_parameters_1, tuning_parameters_2]

        model = LogisticRegression()

        CV=5
        VERBOSE=0

        # This training took 8,5 hours - best model is the one used in the next section
        for i in range(2):
            start = time()
            random_grid_tuned = GridSearchCV(model, tuning_array[i], cv=CV, scoring='f1_micro',
            random_grid_tuned.fit(X_train_transform, Y_train)
            t = time()-start

            # Report result
            # There will be a lot of different combinations and it will take longer for it to se
            b0, m0= FullReport(random_grid_tuned, X_test_transform, Y_test, t)

```

```
sys.stdout = old_stdout
```

```
# OUTPUT:
# 1. Tuning block
# CTOR for best model: LogisticRegression(C=10, class_weight=None,
#           dual=False, fit_intercept=True,
#           intercept_scaling=1, max_iter=1000, multi_class='warn',
#           n_jobs=None, penalty='l2', random_state=None, solver='lbfgs',
#           tol=0.0001, verbose=0, warm_start=False)

# best: data=SignLanguageMnist, score=0.96365,
# model=LogisticRegression(C=10,max_iter=1000,solver='lbfgs')

# 2. Tuning block
# CTOR for best model: LogisticRegression(C=10, class_weight=None,
#           dual=False, fit_intercept=True,
#           intercept_scaling=1, max_iter=100, multi_class='warn',
#           n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
#           tol=0.0001, verbose=0, warm_start=False)

# best: data=SignLanguageMnist, score=0.96361,
# model=LogisticRegression(C=10,max_iter=100,penalty='l1',solver='liblinear')
```

```
In [6]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix
import numpy as np
import matplotlib.pyplot as plt

# Best model from grid search
model = LogisticRegression(solver='lbfgs', multi_class='auto', max_iter=10000,
                           penalty='l2');

# Inspiration from page 102 - 105, Chapter 3, 'Hands-On Machine Learning'
print("Cross Validation Matrix")
score = cross_val_score(model, X_train_transform, Y_train, cv=3, scoring="accuracy")
print(score)

print("Confusion Matrix")
Y_pred = cross_val_predict(model, X_train_transform, Y_train, cv=3)
conf_mx = confusion_matrix(Y_train, Y_pred)
print(conf_mx)

# Plot confusion matrix
plt.matshow(conf_mx, cmap=plt.cm.gray)
```

```

# Plot errors
row_sums = conf_mx.sum(axis=1, keepdims=True)
norm_conf_mx = conf_mx / row_sums

np.fill_diagonal(norm_conf_mx, 0)
plt.matshow(norm_conf_mx, cmap=plt.cm.gray)
plt.show()

```

Cross Validation Matrix

```
[1.          0.99978147 1.          ]
```

Confusion Matrix

```


[[1126    0    0    0    0    0    0    0    0    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0 1009    0    0    0    0    0    0    0    0    0    0    0    0
   0    1    0    0    0    0    0    0    0    0    0]
 [  0    0 1144    0    0    0    0    0    0    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0 1196    0    0    0    0    0    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0  957    0    0    0    0    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0 1204    0    0    0    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0 1090    0    0    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0 1013    0    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0 1162    0    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0 1114    0    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0 1241    0    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0    0 1055    0    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0    0    0 1151    0
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0    0    0    0 1196
   0    0    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0    0    0    0    0
 1088    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0    0    0    0    0
   0 1279    0    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0    0    0    0    0
   0    0 1294    0    0    0    0    0    0    0    0]
 [  0    0    0    0    0    0    0    0    0    0    0    0    0    0
   0    0    0 1199    0    0    0    0    0    0]


```

```

[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 1186 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 1161 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 1 0 0 0 0 0 0 1081 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 1225 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 1164 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 1118]]

```





In []: